

CONCURRENCY

Bùi Tiến Lên

01/01/2020



KHOA CÔNG NGHỆ THÔNG TIN
TRƯỜNG ĐẠI HỌC KHOA HỌC TỰ NHIÊN

Contents



1. **Tasks and Threads**
2. **Passing Arguments**
3. **Returning Results**
4. **Sharing Data**
5. **Waiting for Events**
6. **Communicating Tasks**



Concurrency

Concept 1

Concurrency, the execution of several tasks simultaneously, is widely used

- to improve throughput (by using several processors for a single computation) or
 - to improve responsiveness (by allowing one part of a program to progress while another is waiting for a response)
-
- The C++ standard-library support is primarily aimed at supporting systems-level concurrency rather than directly providing sophisticated higher-level concurrency models.



Tasks and Threads



Tasks vs Threads

Concept 2

- A *task* can be
 - a function
 - a function object
 - a lambda expression
 - A *thread* is the system-level representation of a task in a program.
 - A task to be executed concurrently with other tasks is launched by constructing a thread with the task as its argument.
-
- C++: Using `#include <thread>`

Tasks vs Threads (cont.)



```
void func() {                // function
    cout << "task A\n";
}
class FuncClass {           // function object
public:
    void operator()() {
        cout << "Task B\n";
    }
};
void doSomething() {
    FuncClass funcObj;
    thread t1(func);
    thread t2(funcObj);
    thread t3([]() {cout << "task C\n"; });
    t1.join();                // wait for t1
    t2.join();                // wait for t2
    t3.join();                // wait for t3
}
```

std::thread Members



Member Name	Description
joinable	check if thread joinable
get_id	get ID of thread
native_handle	get native handle for thread
hardware_concurrency	get number of concurrent threads supported by hardware
join	wait for thread to finish executing
detach	permit thread to execute independently
swap	swap threads



The `std::this_thread` Namespace

Name	Description
<code>get_id</code>	get ID of current thread
<code>yield</code>	suggest rescheduling current thread so as to allow other threads to run
<code>sleep_for</code>	blocks execution of current thread for at least specified duration
<code>sleep_until</code>	blocks execution of current thread until specified time reached



Passing Arguments



Passing Arguments

- A task needs data to work upon.

```
void func(vector<double>& v) {...} // function do something with v
    ...
}
class FuncClass {                // function object do something with v
private:
    vector<double>& v;
public:
    FuncClass(vector<double>& vv):v(vv) {}
    void operator()() {...}      // application operator
};
int main() {
    vector<double> some_vec {1,2,3,4,5,6,7,8,9};
    vector<double> other_vec {10,11,12,13,14};
    FuncClass funcObj(other_vec);
    thread t1(f, ref(some_vec)); // executes in a separate thread
    thread t2(funcObj);          // executes in a separate thread
    t1.join();
    t2.join();
}
```



Returning Results



Returning Results

- Pass the input data by const reference
- Pass the location of a place to deposit the result as a separate argument

```
void func(const vector<double>& v, double* res) {  
    *res = accumulate(v.begin(), v.end(), 0);  
}  
void doSomething() {  
    vector<double> v(10000, 1);  
    double res;  
    thread t(func, v, &res);  
    t.join();  
    cout << res << endl;  
}
```



Sharing Data



Problem of Sharing Data

Concept 3

A **race condition** occurs when two or more threads can access shared data and they try to change it at the same time

```
unsigned long long counter = 0;
void func() {
    for (unsigned long long i = 0; i < 1000000; ++i) {
        ++counter;
    }
}
int main() {
    thread t1(func);
    thread t2(func);
    t1.join();
    t2.join();
    cout << counter << endl;
}
```

Mutexes



Concept 4

A **mutex** is a locking mechanism used to synchronize access to a shared resource by providing mutual exclusion.

- A mutex has two basic operations:
 - acquire: `lock` the mutex
 - release: `unlock` the mutex
- A mutex can be held by **only one thread** at any given time.
- C++: Using `#include <mutex>`



Mutexes (cont.)

```
mutex m;  
unsigned long long counter = 0;  
void func() {  
    for (unsigned long long i = 0; i < 1000000; ++i) {  
        m.lock();  
        ++counter;  
        m.unlock();  
    }  
}  
int main() {  
    thread t1(func);  
    thread t2(func);  
    t1.join();  
    t2.join();  
    cout << counter << endl;  
}
```


Deadlock



Concept 5

A **deadlock** is a state in which each member of a group is waiting for another member, including itself, to take action.

Concept 6

A **livelock** is similar to a *deadlock*, except that the states of the processes involved in the livelock constantly change with regard to one another, none progressing.

Deadlock (cont.)



- Consider the following program

```
mutex m;  
void recursiveFunction(int i) {  
    m.lock();  
    if(i > 0) {  
        cout << i;  
        recursiveFunction(i-1);  
    }  
    m.unlock();  
}  
int main() {  
    recursiveFunction(1);  
}
```



Reentrant mutex

Concept 7

A **reentrant mutex** (**recursive mutex**) is a particular type of mutual exclusion mutex that may be locked multiple times by the same process/thread, without causing a **deadlock**.

```
recursive_mutex m;  
void recursiveFunction(int i) {  
    m.lock();  
    if(i > 0) {  
        cout << i;  
        recursiveFunction(i-1);  
    }  
    m.unlock();  
}  
int main() {  
    recursiveFunction(1);  
}
```

Use of resource handles



- Use of RAI resource handles, such as `scoped_lock` and `unique_lock`, is simpler and far safer than explicitly locking and unlocking mutex.

```
mutex m;  
unsigned long long counter = 0;  
void func() {  
    for (unsigned long long i = 0; i < 1000000; ++i) {  
        unique_lock<mutex> lck(m); // acquire mutex  
        ++counter;  
    } // release mutex implicitly  
}  
int main() {  
    thread t1(func);  
    thread t2(func);  
    t1.join();  
    t2.join();  
    cout << counter << endl;  
}
```

Shared mutexes



- The basic mutex allows one thread at a time to access data. One of the most common ways of sharing data is among many readers and a single writer. This “reader-writer lock” idiom is supported by `shared_mutex`.
- A reader will acquire the mutex “shared” so that other readers can still gain access, whereas a writer will demand exclusive access.

```
shared_mutex mx;           // a mutex that can be shared
void reader() {
    shared_lock<shared_mutex> lck(mx); // willing to share access
                                     with other
                                     // readers
    // ... read ...
}
void writer() {
    unique_lock<shared_mutex> lck(mx); // needs exclusive (unique)
                                     access
    // ... write ...
}
```



Waiting for Events



Waiting for Events

- Sometimes, a thread needs to wait for some kind of external event, such as another thread completing a task or a certain amount of time having passed.
- The basic support for communicating using external events is provided by `condition_variable`.
- A `condition_variable` is a mechanism allowing one thread to wait for another.



condition_variable Members

Name	Description
<code>notify_one</code>	notify one waiting thread
<code>notify_all</code>	notify all waiting threads
<code>wait</code>	blocks current thread until notified
<code>wait_for</code>	blocks current thread until notified or specified duration passed
<code>wait_until</code>	blocks current thread until notified or specified time point reached



Example

- Consider the classical example of two threads communicating by passing messages through a queue.

```
class Message {      // object to be communicated
    // ...
};
queue<Message> mqueue;    // the queue of messages
condition_variable mcond; // the variable communicating events
mutex mmutex;           // for synchronizing access to mcond
```



Example (cont.)

```
void consumer() {
    while (true) {
        unique_lock<mutex> lck(mmutex); // acquire mmutex
        mcond.wait(lck, [] { return !mqueue.empty(); });
        // release lck and wait;
        // re-acquire lck upon wakeup
        // don't wake up unless mqueue is non-empty

        auto m = mqueue.front(); // get the message
        mqueue.pop();
        lck.unlock();             // release lck (optional)
        // ... process m ...
    }
}
```



Example (cont.)

```
void producer() {  
    while (true) {  
        Message m;  
        // ... fill the message ...  
        unique_lock<mutex> lck(mmutex); // protect operations  
        mqueue.push(m);  
        mcond.notify_one();           // notify  
    } // release lock (at end of scope)  
}
```



Communicating Tasks

Communicating Tasks



- The standard library provides a few facilities to allow programmers to operate at the conceptual level of tasks (work to potentially be done concurrently) rather than directly at the lower level of threads and locks:
 - `future` and `promise` for returning a value from a task spawned on a separate thread
 - `async()` for launching of a task in a manner very similar to calling a function
- C++: `#include <future>`



promise and future

- The important point about future and promise is that they enable a transfer of a value between two tasks without explicit use of a lock
- promise Members

Name	Description
swap	swap two promise objects
<code>get_future</code>	get future associated with promised result
<code>set_value</code>	set result to specified value
<code>set_value_at_thread_exit</code>	set result to specified value while delivering notification only at thread exit
<code>set_exception</code>	set result to specified exception
<code>set_exception_at_thread_exit</code>	set result to specified exception while delivering notification only at thread exit



promise and future (cont.)

- future Members

Name	Description
share	transfer shared state to shared_future object
get	get result
valid	check if future object refers to shared state
wait	wait for result to become available
wait_for	wait for result to become available or time duration to expire
wait_until	wait for result to become available or time point to be reached



Example

```
void f(promise<X>& px) {  
    // a task: place the result in px  
    // ...  
    try {  
        X res;  
        // ... compute a value for res ...  
        px.set_value(res);  
    }  
    catch (...) {  
        // pass the exception to the future's thread  
        px.set_exception(current_exception());  
    }  
}
```




Example (cont.)

```
void g(future<X>& fx) {  
    // a task: get the result from fx  
    // ...  
    try {  
        X v = fx.get();  
        // if necessary, wait for the value to get computed  
        // ... use v ...  
    }  
    catch (...) {  
        // ... handle error ...  
    }  
}
```



async()

- To launch tasks to potentially run asynchronously, we can use `async()`

```
double comp4(vector<double>& v) {  
    // spawn many tasks if v is large enough  
    if (v.size() < 10000) // is it worth using concurrency?  
        return accum(v.begin(), v.end(), 0.0);  
  
    auto v0 = &v[0];  
    auto sz = v.size();  
  
    auto f0 = async(accum, v0, v0+sz/4, 0.0);           // first quarter  
    auto f1 = async(accum, v0+sz/4, v0+sz/2, 0.0);      // second quarter  
    auto f2 = async(accum, v0+sz/2, v0+sz*3/4, 0.0);    // third quarter  
    auto f3 = async(accum, v0+sz*3/4, v0+sz, 0.0);      // fourth quarter  
  
    // collect and combine the results  
    return f0.get()+f1.get()+f2.get()+f3.get();  
}
```

References



Deitel, P. (2016).

C++: How to program.

Pearson.



Gaddis, T. (2014).

Starting Out with C++ from Control Structures to Objects.

Addison-Wesley Professional, 8th edition.



Jones, B. (2014).

Sams teach yourself C++ in one hour a day.

Sams.